



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

Experiment - 5

Student Name: Harjit Singh

UID: 23BCS10849

Branch: BE-CSE

Section/Group: KRG_2B

Semester: 5th

Date of Performance: 22/9/25

Subject Name: Advanced Database and Management System

Subject Code: 23CSP-333

1. Problem Description/Aim:

Medium-Problem Title: Generate 1 million records per ID in ‘transaction_data’ using generate_series() and random() ,create a normal view and a materialized view ‘sales_summary’ with aggregated metrics (total_quantity_sold , total_sales, total_orders) , and compare their performance and execution time.

Procedure (Step-by-Step):

1. Create a large dataset:
 - Create a table names transaction_data (id , value) with 1 million records.
 - take id 1 and 2, and for each id, generate 1 million records in value column
 - Use Generate_series () and random() to populate the data.
2. Create a normal view and materialized view to for sales_summary, which includes total_quantity_sold, total_sales, and total_orders with aggregation.
3. Compare the performance and execution time of both.

Sample Output Description:

The transaction_data table has 2 million rows (1 million per ID) with random values. The normal view sales_summary computes aggregates on the fly, while the materialized view sales_summary_mv stores precomputed results. Queries on the materialized view are much faster, but it needs refreshing when data changes, whereas the normal view always shows up-to-date results.

Hard-Problem Title: Create restricted views in the sales database to provide summarized, non-sensitive data to the reporting team, and control access using DCL commands(GRANT and REVOKE).

Procedure (Step-by-Step):

1. Create restricted views-
 - Define views that show only **aggregated sales data** (e.g., total_sales, total_orders) without exposing sensitive columns like customer details or payment info.



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

2. Assign access to reporting team(or client)-
-Use “GRANT SELECT ON view_name TO reporting_user;” to give access.
3. Revoke access if needed.
-Use “REVOKE SELECT ON view_name FROM reporting_user;” to remove access.
4. Verify access
- Reporting users can query the view but cannot access base tables directly, ensuring security.

Sample Output Description:

The result shows the restricted view providing summarized sales data only like
- Columns shown are - product_id, total_quantity_sold, total_sales, total_orders
- Columns hidden are - Customer names, addresses, payment details

A reporting user querying the view sees something like :

- Product 101 - 5000 units sold, total sales Rs. 12,50,000,500 orders.
- Product 102 - 3200 units sold, total sales Rs. 8,60,000,320 orders.

When the user tries to query the base “sales_transactions” table directly, access is denied, enforcing security.

2. **Objective:** To design and implement secure, efficient data access mechanisms by creating large-scale transaction datasets, summarizing them through normal and materialized views for performance comparison, and enforcing restricted access to sensitive data using views and DCL commands.

3. SQL QUERY AND OUTPUTS -

MEDIUM LEVEL PROBLEM-----

```
Create table TRANSACTION_DATA(id int, val decimal);
INSERT INTO TRANSACTION_DATA(ID,VAL)
SELECT 1,RANDOM()
FROM GENERATE_SERIES(1,1000000);
```

```
INSERT INTO TRANSACTION_DATA(ID,VAL)
SELECT 2,RANDOM()
FROM GENERATE_SERIES(1,1000000);
SELECT * FROM TRANSACTION_DATA;
```



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

CREATE or REPLACE VIEW SALES_SUMMARY AS

SELECT

ID,
COUNT(*) AS total_quantity_sold,
sum(val) AS total_sales,
count(distinct id) AS total_orders
FROM TRANSACTION_DATA
GROUP BY ID;

EXPLAIN ANALYZE

SELECT * FROM SALES_SUMMARY;

CREATE MATERIALIZED VIEW SALES_SUMM AS

SELECT

ID,
COUNT(*) AS total_quantity_sold,
sum(val) AS total_sales,
count(distinct id) AS total_orders
FROM TRANSACTION_DATA
GROUP BY ID;

EXPLAIN ANALYZE

SELECT * FROM SALES_SUMM;

```
6   INSERT INTO TRANSACTION_DATA(ID,VAL)
7   SELECT 2,random()
8   FROM generate_series(1,1000000);
9   SELECT * FROM TRANSACTION_DATA;
```

Data Output Messages Notifications

	id integer	val numeric
1	1	0.748060017288284
2	1	0.158813530918857
3	1	0.482094772953915
4	1	0.461220286286965
5	1	0.601375928005661
6	1	0.120882758237791
7	1	0.626445464971291
8	1	0.448741750697511
9	1	0.127332205463045



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

```
21   SELECT * FROM SALES_SUMMARY; /*Simple view */
```

Data Output Messages Notifications

A screenshot of a database query results table. The table has four columns: id, total_quantity_sold, total_sales, and total_orders. The data shows two rows: one with id 1, total_quantity_sold 2000000, total_sales 1000226.201610874170319933640, and total_orders 1; and another with id 2, total_quantity_sold 1000000, total_sales 499473.47586932728250459408, and total_orders 1.

	id integer	total_quantity_sold bigint	total_sales numeric	total_orders bigint
1	1	2000000	1000226.201610874170319933640	1
2	2	1000000	499473.47586932728250459408	1

```
20   EXPLAIN ANALYZE
```

```
21   SELECT * FROM SALES_SUMMARY; /*Simple view */
```

Data Output Messages Notifications

A screenshot of an Explain Analyze query plan table. The table has a single column labeled "QUERY PLAN". It contains eight numbered rows detailing the execution plan: 1. GroupAggregate (cost=471514.97..509014.99 rows=2 width=52) (a), 2. Group Key: transaction_data.id, 3. -> Sort (cost=471514.97..479014.97 rows=3000000 width=15) (a), 4. Sort Key: transaction_data.id, 5. Sort Method: external merge Disk: 73504kB, 6. -> Seq Scan on transaction_data (cost=0.00..46224.00 rows=3000000 width=15), 7. Planning Time: 0.135 ms, 8. Execution Time: 4396.880 ms.

QUERY PLAN
text
1 GroupAggregate (cost=471514.97..509014.99 rows=2 width=52) (a)
2 Group Key: transaction_data.id
3 -> Sort (cost=471514.97..479014.97 rows=3000000 width=15) (a)
4 Sort Key: transaction_data.id
5 Sort Method: external merge Disk: 73504kB
6 -> Seq Scan on transaction_data (cost=0.00..46224.00 rows=3000000 width=15)
7 Planning Time: 0.135 ms
8 Execution Time: 4396.880 ms

```
33   SELECT * FROM SALES_SUMM; /*Materialized view*/
```

Data Output Messages Notifications

A screenshot of a materialized view query results table. The table has four columns: id, total_quantity_sold, total_sales, and total_orders. The data shows two rows: one with id 1, total_quantity_sold 1000000, total_sales 500106.667545326356598143529, and total_orders 1; and another with id 2, total_quantity_sold 1000000, total_sales 499473.47586932728250459408, and total_orders 1.

	id integer	total_quantity_sold bigint	total_sales numeric	total_orders bigint
1	1	1000000	500106.667545326356598143529	1
2	2	1000000	499473.47586932728250459408	1



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

```
32 | EXPLAIN ANALYZE
33  SELECT * FROM SALES_SUMM; /*Materialized view*/
```

Data Output Messages Notifications

Showing rows: 1

QUERY PLAN	
text	
1	Seq Scan on sales_summ (cost=0.00..20.20 rows=1020 width=52) (actual time=0.017..0.018 rows=2 loops=...)
2	Planning Time: 0.063 ms
3	Execution Time: 0.032 ms

1	Seq Scan on sales_summ (cost=0.00..20.20 rows=1020 width=52) (actual time=0.017..0.018 rows=2 loops=...)
2	Planning Time: 0.063 ms
3	Execution Time: 0.032 ms

OUTPUT -

As we can see that the execution time using the materialized view is very less as compared to the simple view's execution time.

HARD PROBLEM -----

```
CREATE TABLE customer_data (
    transaction_id SERIAL PRIMARY KEY,
    customer_name VARCHAR(100),
    email VARCHAR(100),
    phone VARCHAR(15),
    payment_info VARCHAR(50), -- sensitive
    order_value DECIMAL,
    order_date DATE DEFAULT CURRENT_DATE
);
```

-- Insert sample data

```
INSERT INTO customer_data (customer_name, email, phone, payment_info, order_value)
VALUES
('Mandeep Kaur', 'mandeep@example.com', '9040122324', '1234-5678-9012-3456', 500),
('Mandeep Kaur', 'mandeep@example.com', '9040122324', '1234-5678-9012-3456', 1000),
('Jaskaran Singh', 'jaskaran@example.com', '9876543210', '9876-5432-1098-7654', 700),
('Jaskaran Singh', 'jaskaran@example.com', '9876543210', '9876-5432-1098-7654', 300);
```



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

```
CREATE OR REPLACE VIEW RESTRICTED_SALES_DATA AS
```

```
SELECT
```

```
CUSTOMER_NAME,
```

```
COUNT(*) AS total_orders,
```

```
SUM(order_value) as total_sales
```

```
from customer_data
```

```
group by customer_name;
```

```
select * from restricted_sales_data;
```

```
CREATE USER CLIENT1 WITH PASSWORD 'REPORT1234';
```

```
GRANT SELECT ON RESTRICTED_SALES_DATA TO CLIENT1;
```

```
REVOKE SELECT ON RESTRICTED_SALES_DATA FROM CLIENT1;
```

The screenshot shows a pgAdmin 4 interface. The top bar displays the session as "Mandeep/client1@PostgreSQL 17". A message bar indicates "The session is idle and there is no current transaction." Below the toolbar, there are tabs for "Query" (which is selected) and "Query History". The main query editor area contains the following code:

```
62 group by customer_name;
63
64 select * from restricted_sales_data;
65
```

Below the editor, there are tabs for "Data Output", "Messages" (which is selected), and "Notifications". The "Messages" tab displays the error output:

```
ERROR: permission denied for view restricted_sales_data
SQL state: 42501
```



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

Mandeep/postgres@PostgreSQL 17

No limit

Query History

```
65
66 CREATE USER CLIENT1 WITH PASSWORD 'REPORT1234';
67 GRANT SELECT ON RESTRICTED_SALES_DATA TO CLIENT1;
68 REVOKE SELECT ON RESTRICTED_SALES_DATA FROM CLIENT;
```

Data Output Messages Notifications

GRANT

Query returned successfully in 154 msec.

Mandeep/client1@PostgreSQL 17

No limit

Query History

```
62 group by customer_name;
63
64 select * from restricted_sales_data;
65
```

Data Output Messages Notifications

SQL

	customer_name character varying (100)	total_orders bigint	total_sales numeric
1	Jaskaran Singh	2	1000
2	Mandeep Kaur	2	1500



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

Mandeep/postgres@PostgreSQL 17

No limit

Query History

```
64 select * from restricted_sales_data;
65
66 CREATE USER CLIENT1 WITH PASSWORD 'REPORT1234';
67 GRANT SELECT ON RESTRICTED_SALES_DATA TO CLIENT1;
68 REVOKE SELECT ON RESTRICTED_SALES_DATA FROM CLIENT1;
```

Data Output Messages Notifications

REVOKE

Query returned successfully in 163 msec.

Mandeep/client1@PostgreSQL 17

No limit

Query History

```
63
64 select * from restricted_sales_data;
65
66 CREATE USER CLIENT1 WITH PASSWORD 'REPORT1234';
67 GRANT SELECT ON RESTRICTED_SALES_DATA TO CLIENT1;
68 REVOKE SELECT ON RESTRICTED_SALES_DATA FROM CLIENT1;
```

Data Output Messages Notifications

ERROR: permission denied for view restricted_sales_data

SQL state: 42501